



Practical Implementation

1. It is enough to update $A(k + 1 : n, k + 1 : n)$
2. We can overwrite $A(k + 1 : n, k)$ with $L(k + 1 : n, k)$



(2) Since we know that entries below the k^{th} col $A(k+1:n, k)$ is going to be zero, we use this to store the multipliers: $L(k+1:n, k)$



LU Algorithm

```
for  $k = 1 : n - 1$  do
  (*)  $A(k + 1 : n, k) = A(k + 1 : n, k) / A(k, k)$  ← compute multipliers
  for  $i = k + 1 : n$  do
    for  $j = k + 1 : n$  do
       $A(i, j) = A(i, j) - \underbrace{A(i, k)A(k, j)}$  ← update the rows
    end for
  end for
end for
```

↑ These are multipliers stored in (*) above

Vectorize j th loop.

LU Algorithm: After Vectorization of j th loop

for $k = 1 : n - 1$ **do**

$$A(k + 1 : n, k) = A(k + 1 : n, k) / A(k, k)$$

for $i = k + 1 : n$ **do**

$$\underbrace{A(i, k + 1 : n)}_{R_i} = \underbrace{A(i, k + 1 : n)}_{R_i} - \underbrace{A(i, k)}_{m_i} \underbrace{A(k, k + 1 : n)}_{R_k}$$

end for

end for

$$R_i \leftarrow R_i - m_i R_k$$

After vectorizing j -loop, we see row operations!

Parallel Computing: MPI

Moore's Law

Microprocessor Transistor Counts 1971-2011 & Moore's Law

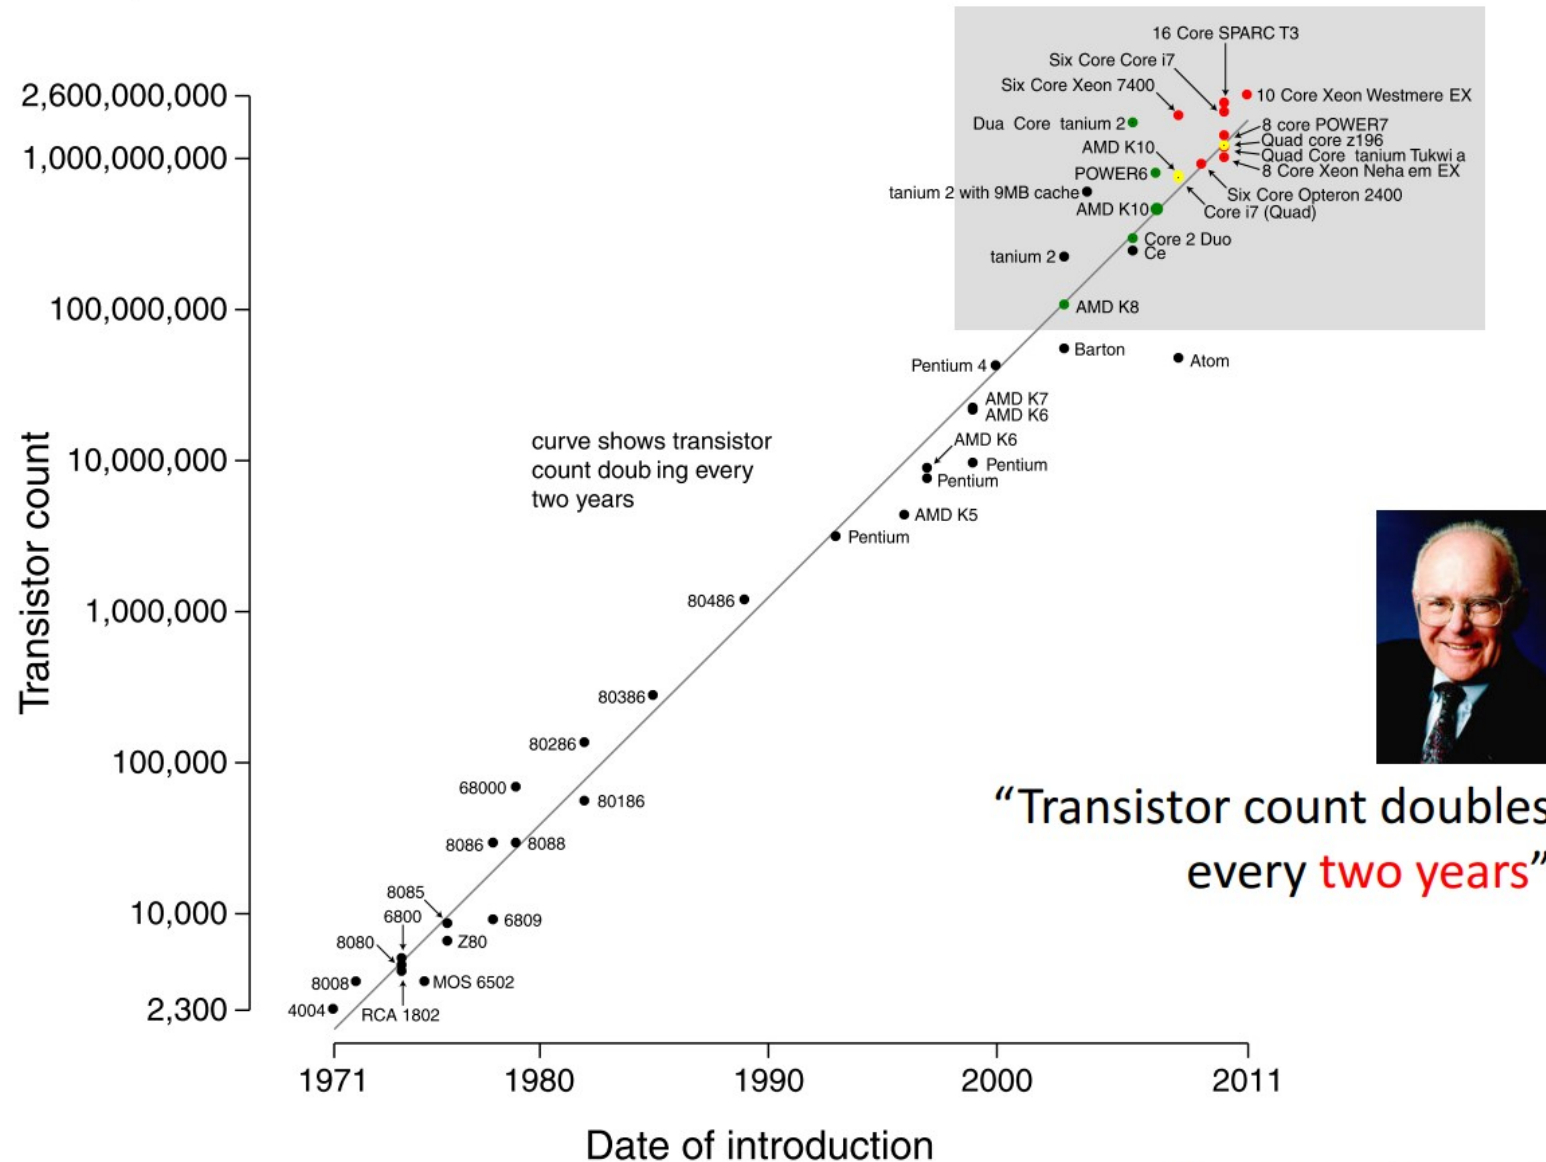



Illustration from Wikipedia

Moore's Law



Top 500 Supercomputers





Image taken from www.top500.org

Operating Systems on HPC systems



HPC systems
are mostly
dominated by
Unix-like system

Motivation for Parallel Computing

1. Want to run the same program faster


2. Want to run bigger datasets

Datasets may be so big that it does not fit in RAM

3. Want to reduce financial cost or power consumption

PhD students waiting for weeks for results!

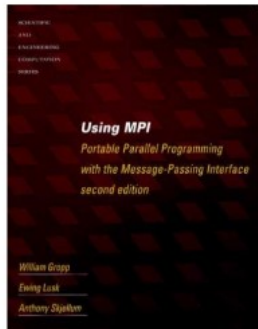

Distributed Memory Architecture



Message Passing Interface

1. MPI: Library specification, not an implementation
2. Most important implementations: OpenMPI, Intel MPI
3. Specific routines for:
 - 3.1 Point-to-Point Communication
 - 3.2 Collective communications
 - 3.3 Topology setup
 - 3.4 Parallel I/O
4. Binding for C/C++

References for MPI



- **MPI standards:** <http://www.mpi-forum.org/docs/>
- **MPI: The Complete Reference** (M. Snir, S. Otto, S. Huss-Lederman, D. Walker, J. Dongarra)
Available from <http://switzernet.com/people/emin-gabrielyan/060708-thesis-ref/papers/Snir96.pdf>
- **Using MPI: Portable Parallel Programming with the Message Passing Interface**, 2nd ed. (W. Gropp, E. Lusk, A. Skjellum).

History of MPI

Started in 1992 (Workshop for standards for Message Passing in a Distributed Memory Environment) with support from vendors, library writers, and academia.

1. MPI Version 1.0 (May 1994)
 - 1.1 Final pre draft in 1993
 - 1.2 Final version June 1994
2. MPI version 2.0
 - 2.1 Support for one sided communication
 - 2.2 Support for process management
 - 2.3 Support for parallel I/O
3. MPI version 3.0
 - 3.1 Support for non-blocking collective routine
 - 3.2 New one sided communication routines

Hello World in MPI

```
#include <mpi.h>
#include <iostream>
#include <cstdlib>

using namespace std;

int main(int argc, char* argv[]) {
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );

    cout << "Hello World from process" << rank << "/" << size << endl;

    MPI_Finalize();
    return EXIT_SUCCESS;
}
```

Output **order** is
random


```
john@doe ~]$ mpirun -np 4 ./helloWorld
Hello World from process 2/4
Hello World from process 3/4
Hello World from process 0/4
Hello World from process 1/4
```

Basic MPI routines

- `int MPI_Init(int *argc, char ***argv)`
 - Initialization: All process must call this prior to any MPI routine
 - Strips off possible arguments provided by “mpirun”
- `int MPI_Finalize(void)`
 - Cleanup: all processes must call this routine at the end of the program
 - All pending communication should have finished before calling this


Basic MPI Routines

- `int MPI_Comm_size(MPI_Comm comm, int *size);`
 - Returns the size of the “Communicator” associated with “comm”
 - Communicator is user defined subset of processes
 - `MPI_COMM_WORLD` = communicator that involves all processes



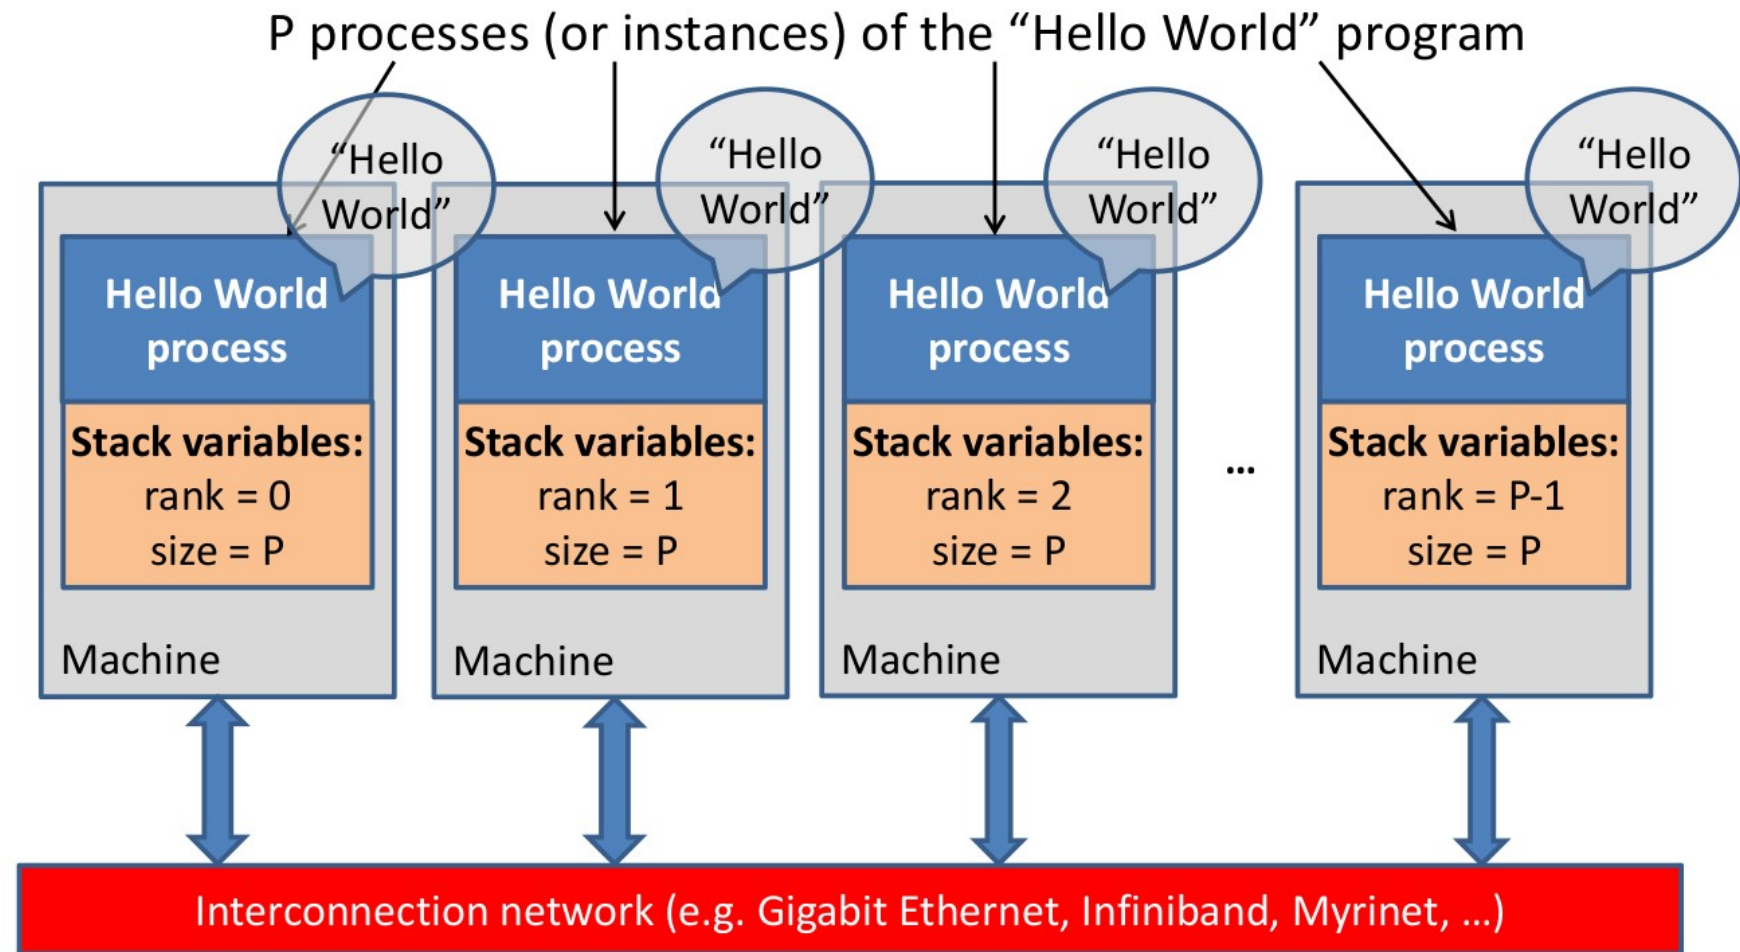
Basic MPI Routines

- `int MPI_Comm_rank(MPI_Comm comm, int *rank);`
 - Return the rank of the process in the Communicator
 - Range: $[0 \dots \text{size} - 1]$



After calling this,
unique id called
rank is assigned
to each CPUs.

MPI Mechanisms

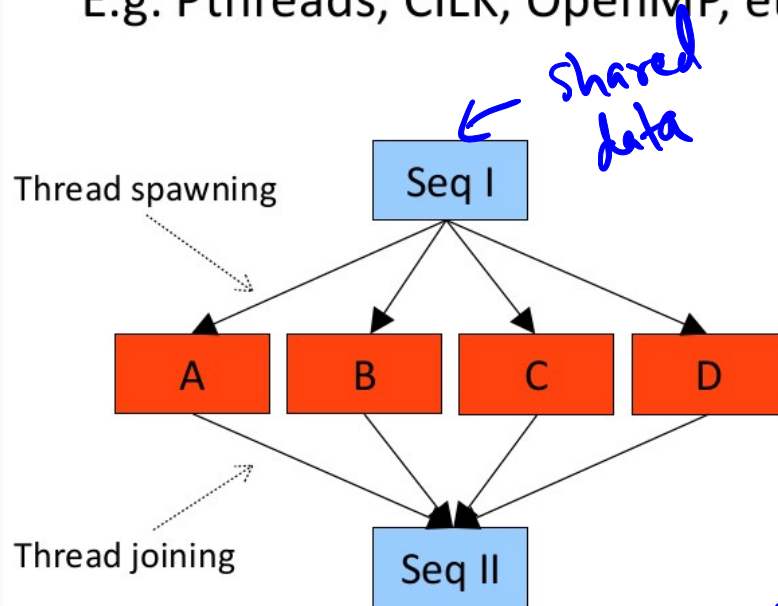


- MPI launches P independent processes: each process is an instance of same program

Multithreading Versus Multiprocessing

Multi-threading

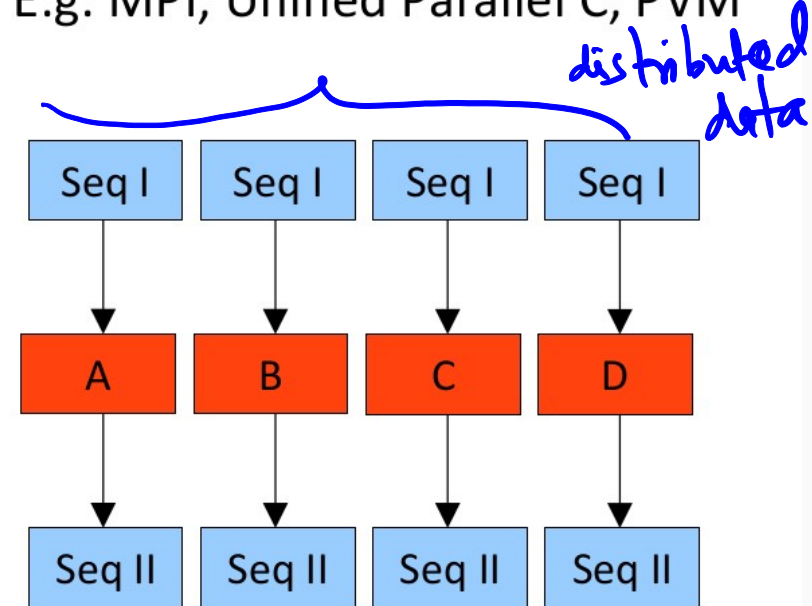
Single process
Shared memory address space
Protect data against simultaneous writing
Limited to a single machine
E.g. Pthreads, CILK, OpenMP, etc.



↑ Threads work on shared data

Message passing

Multiple processes
Separate memory address spaces
Explicitly communicate everything
Multiple machines possible
E.g. MPI, Unified Parallel C, PVM



↑ Processes work on their local data

Compiling and Running MPI programs

- Compiling: `mpicc -O3 main.cpp -o main`
- Running: `mpirun -np <number of program instances> <your program>`

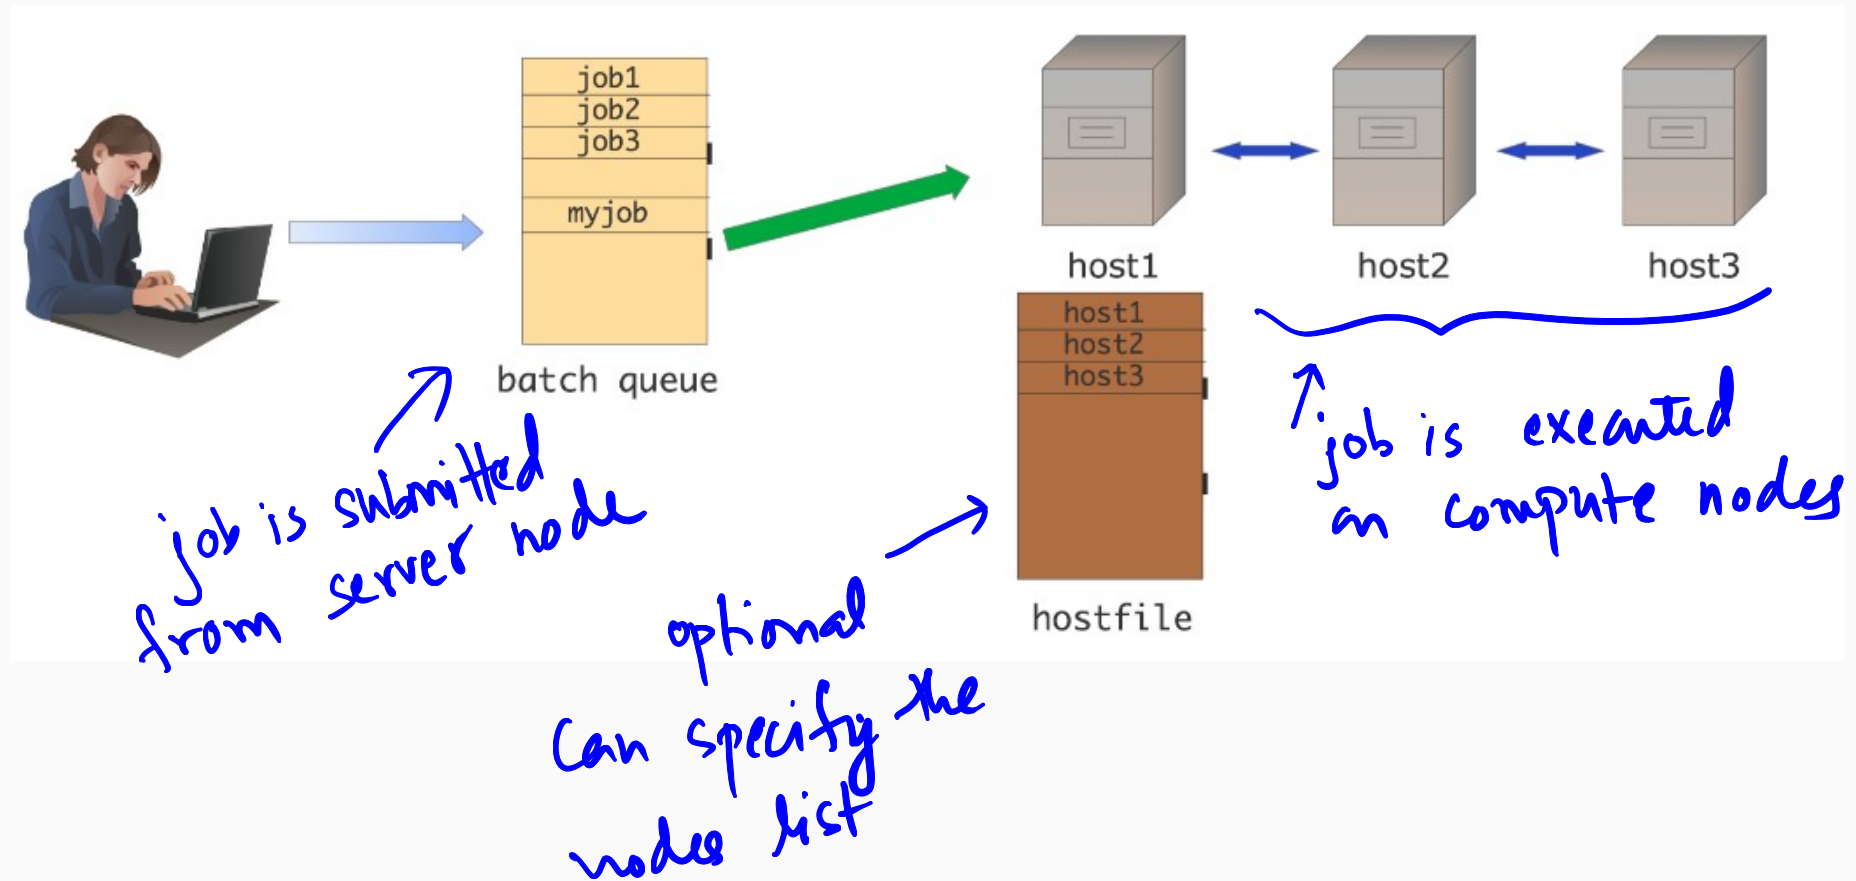
Using MPI on ABACUS cluster at IIITH

1. check modules: `module avail`
2. load modules: `module load intel/2017a`
3. submit job using job script: `sbatch job_script.sh`

More details:

http://hpc.iiit.ac.in/wiki/index.php/Main_Page

Batch Job Submission in MPI



MPI: Point-to-Point Communication

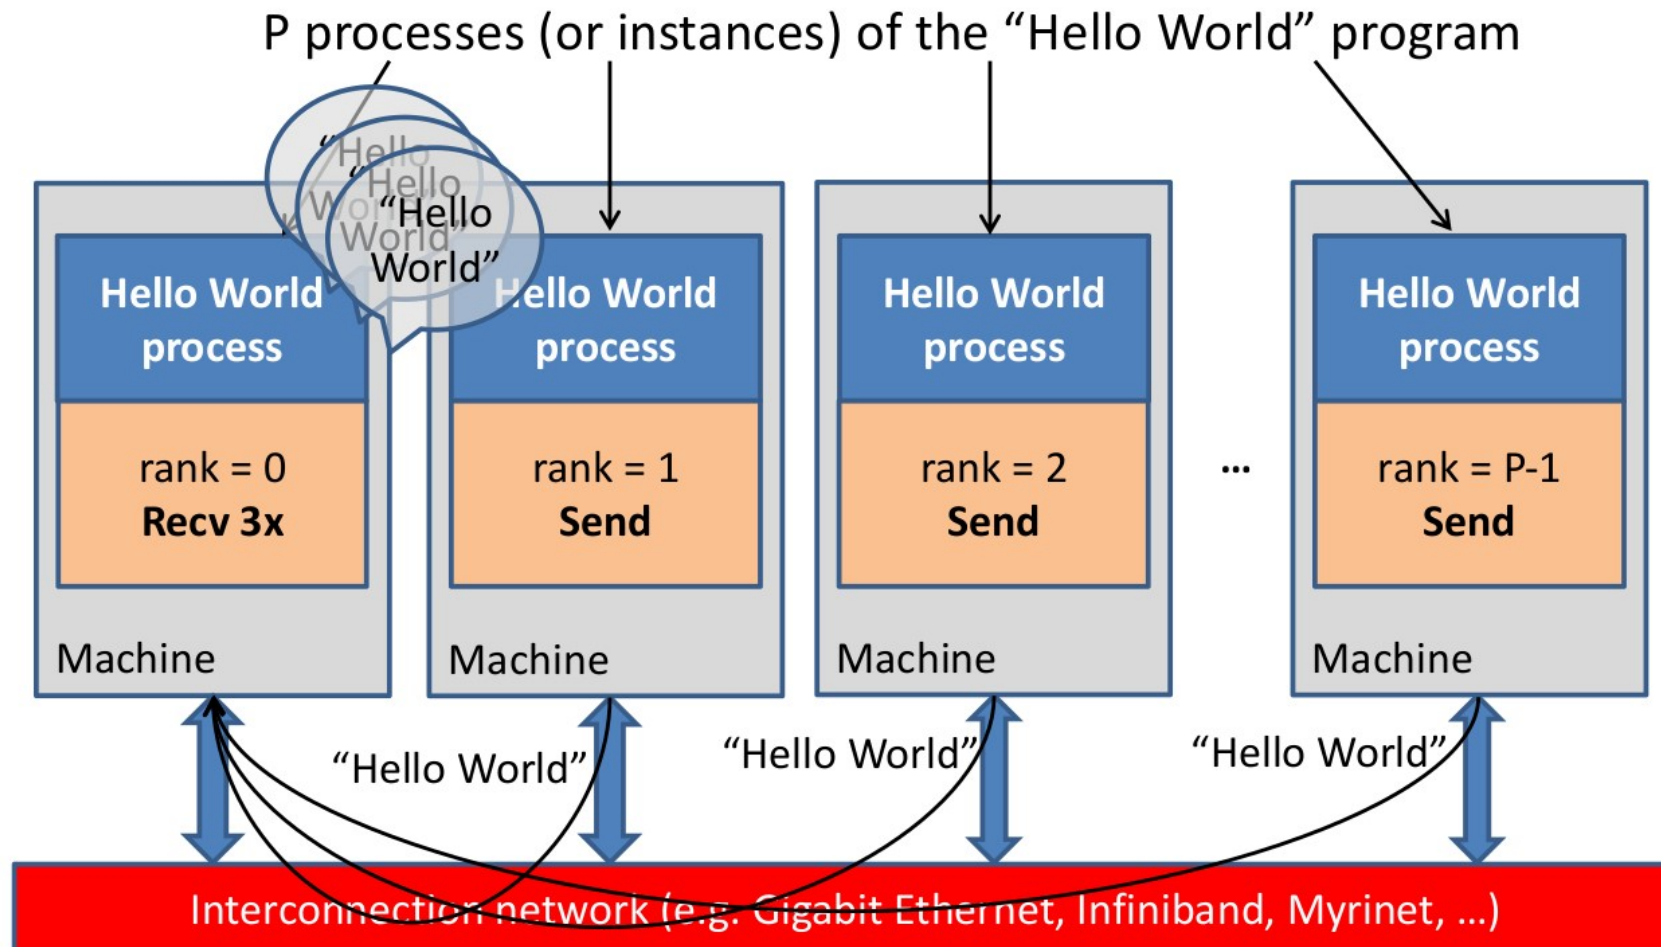
```
...  
int rank, size, count;  
char b[40];  
MPI_Status status;  
  
... // init MPI and rank and size variables  
  
if (rank != 0) {  
    char * str = "Hello World";  
    MPI_Send(str, 12, MPI_CHAR, 0, 123, MPI_COMM_WORLD);  
} else {  
    for (int i = 1; i < size; i++) {  
        MPI_Recv(b, 40, MPI_CHAR, i, MPI_ANY_TAG, MPI_COMM_WORLD, &status);  
        MPI_Get_count(&status, MPI_CHAR, &count);  
        printf("I received %s from process %d with size %d and tag %d\n",  
              b, status.MPI_SOURCE, count, status.MPI_TAG);  
    }  
}  
...
```

branching on rank

```
john@doe ~]$ mpirun -np 4 ./ptpcomm
```

```
I received Hello World from process 1 with size 12 and tag 123  
I received Hello World from process 2 with size 12 and tag 123  
I received Hello World from process 3 with size 12 and tag 123
```

MPI: Point-to-Point Communication View



MPI: send and receive

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype,  
int dest, int tag, MPI_Comm comm)
```

- **buf**: pointer to the message to send
- **count**: number of items to send
- **datatype**: datatype of each item
- **dest**: rank of destination process
- **tag**: value to identify the message [0 ... at least (32 767)]
- **comm**: communicator specification (e.g. MPI_COMM_WORLD)

MPI: send and receive

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,  
int source, int tag, MPI_Comm comm, MPI_Status *status)
```

- **buf:** pointer to the buffer to store received data
- **count:** upper bound (!) of the number of items to receive
- **datatype:** datatype of each item
- **source:** rank of source process (or MPI_ANY_SOURCE)
- **tag:** value to identify the message (or MPI_ANY_TAG)
- **comm:** communicator specification (e.g. MPI_COMM_WORLD)
- **status:** structure that contains MPI_SOURCE, MPI_TAG, MPI_ERROR

Sending and Receiving

Two-sided communication:

- Both the sender and receiver are involved in data transfer
- Posted send must match receive

When do MPI_Send and MPI_recv match ?

- Rank of receiver process
- Rank of sending process
- Tag (custom value to distinguish messages from same sender)
- Communicator

Rationale for communicators:

- Used to create subsets of processes